

RDD Execution Steps

- 1. **Data Loading** – Read data from HDFS, S3, or databases.
- 2. **Partitioning** – Data is split into **smaller chunks (partitions)**.
- 3. **Execution** – Spark processes each partition **in parallel**.

RDD Properties

| Property | Description |
|------------------|---|
| Immutable | Cannot be modified after creation. |
| Partitioned | Distributed across nodes for parallel processing. |
| Fault-Tolerant | Can recover data from lineage. |
| Lazily Evaluated | Execution is delayed until an action is called. |

1. Lazy Evaluation in Spark

Why Lazy Evaluation?

- Optimizes execution by **building a DAG (Directed Acyclic Graph)**.
- **Only executes transformations** when an action is triggered.

Transformations (Lazy) vs Actions (Triggers Execution)

| Operation Type | Description | Examples |
|-----------------|------------------------------|---|
| Transformations | Creates a new RDD | <code>map()</code> , <code>filter()</code> , <code>flatMap()</code> |
| Actions | Executes all transformations | <code>collect()</code> , <code>count()</code> , <code>reduce()</code> |

2. Transformations & Actions in RDD

Narrow Transformations (No Shuffling)

- Each output partition **depends on a single input partition**.
- **Faster execution.**
- **Examples:** `map()` , `filter()` , `flatMap()` .

Wide Transformations (Shuffling Required)

- Data is **redistributed across partitions**.
- **Slower due to network communication.**
- **Examples:** `reduceByKey()` , `groupByKey()` , `sortByKey()` .

Example

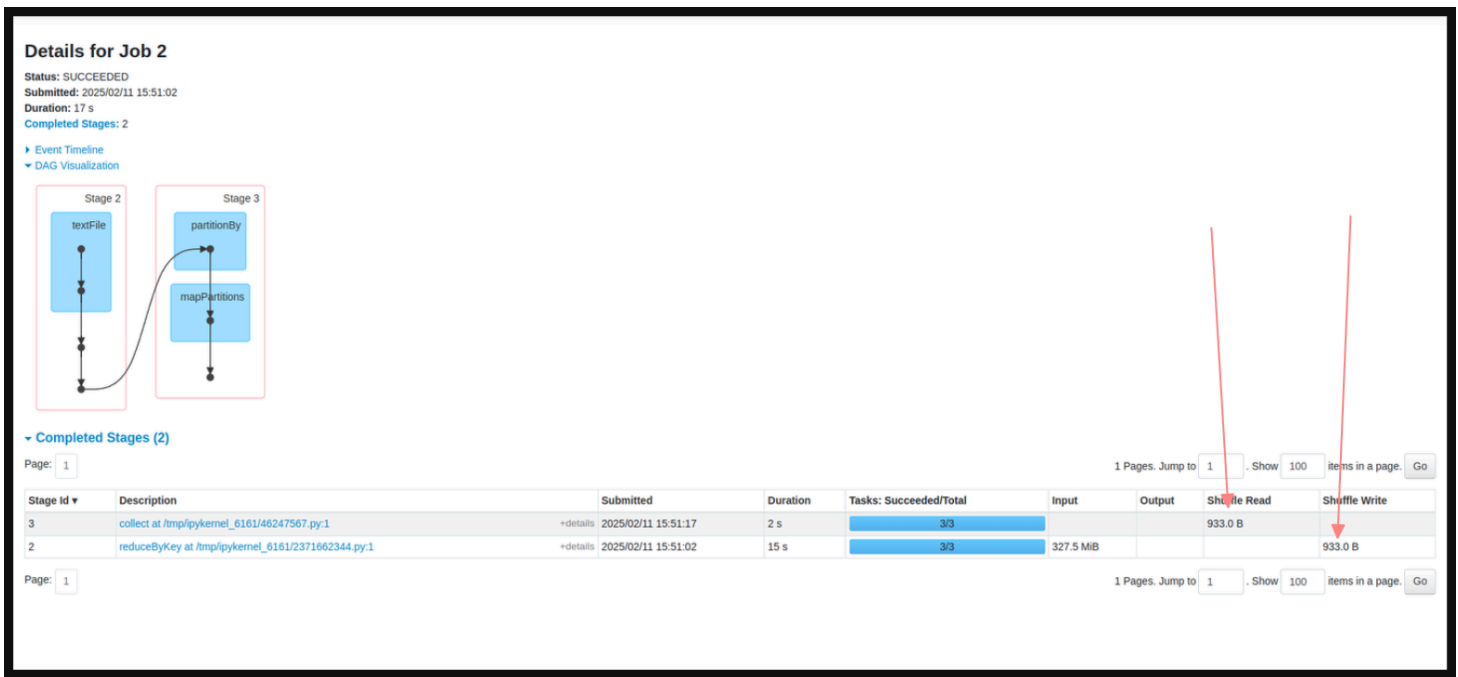
ReduceByKey vs GroupByKey in Spark

1. ReduceByKey

- **Type:** Wide Transformation
- **Definition:**
 - Aggregates values for each key using a **specified associative and commutative reduce function**.
 - **Combiner Optimization:** Combines the values **locally within each partition** before shuffling data across the network.
 - **Result:** Produces a single output value per key.
- **Example Function:**
 - `lambda x, y: x + y`
- **Key Points:**
 - Reduces the amount of data shuffled over the network.
 - More efficient for large datasets.
- **Output:**
 - RDD of `(key, aggregated_value)` pairs.
- **Example:**

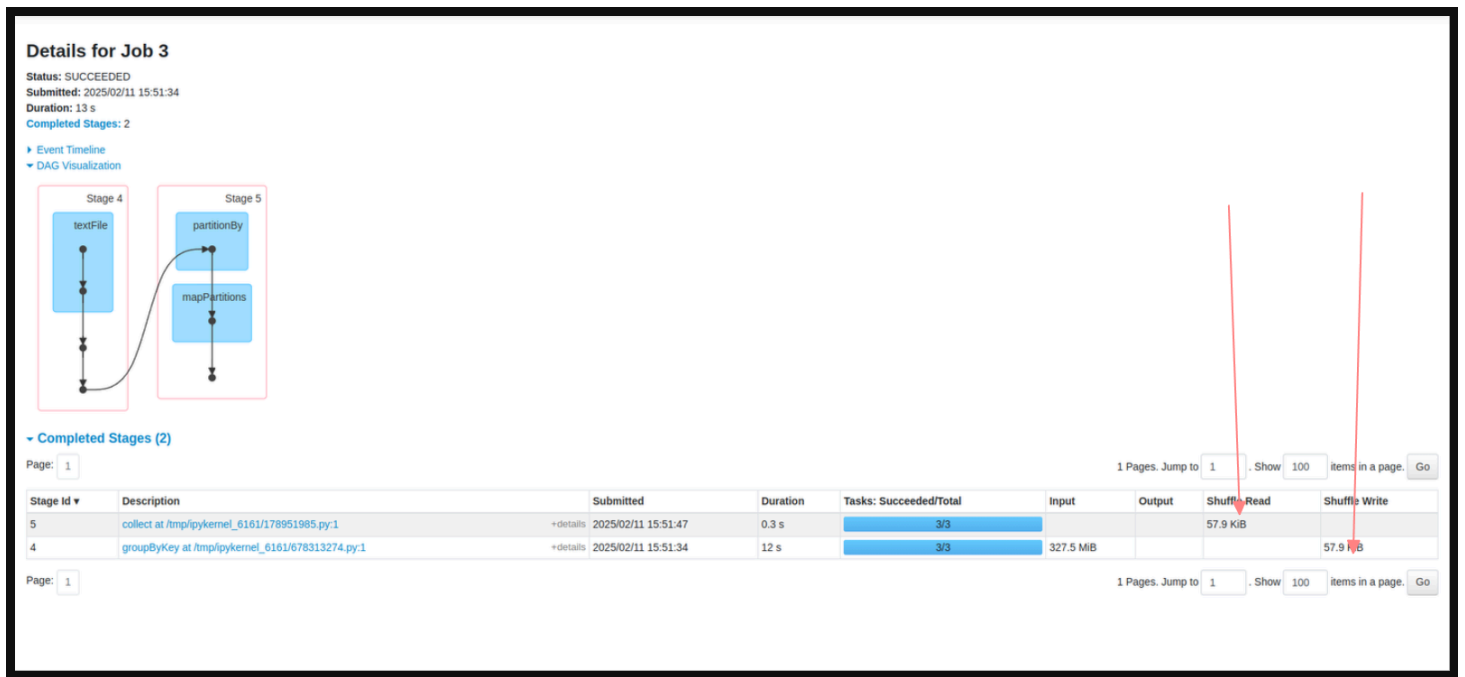
Input: `[(a, 1), (b, 2), (a, 3), (b, 4)]`

Output: `[(a, 4), (b, 6)]`



2. GroupByKey

- **Type:** Wide Transformation
- **Definition:**
 - Groups all values associated with each key into a **single iterable collection**.
 - **No Aggregation:** Only groups data, does not perform aggregation.
- **Key Points:**
 - Shuffles all data across the network, which can be expensive for large datasets.
 - Suitable when aggregation is not required and all values per key are needed.
- **Output:**
 - RDD of (key, Iterable[values]) pairs.
- **Example:**
Input: [(a, 1), (b, 2), (a, 3), (b, 4)]
Output: [(a, [1, 3]), (b, [2, 4])]



Differences Between ReduceByKey and GroupByKey

| Feature | ReduceByKey | GroupByKey |
|----------------|---|--|
| Operation Type | Aggregates values per key. | Groups values into an iterable per key. |
| Shuffling | Combines values locally before shuffling. | Shuffles all data directly. |
| Efficiency | More efficient for large datasets. | Less efficient due to higher shuffle cost. |
| Use Case | When aggregation is needed. | When all values for a key are required. |
| Output | (key, aggregated_value) pairs. | (key, Iterable[values]) pairs. |

Shuffle in Wide Transformations

- **Shuffle** involves **redistributing data** across the cluster.
- It is triggered when:
 - Data needs to be grouped or aggregated across partitions.
- Shuffling is **resource-intensive** and **time-consuming**, as it involves:
 - Writing intermediate data to disk.
 - Network communication between nodes.

Optimization Tip

- Try to minimize **wide transformations** (e.g., shuffles) to improve performance.
- Use **narrow transformations** whenever possible to avoid the overhead of data shuffling.